



Katholieke
Universiteit
Leuven

Department of
Computer Science

DOCUMENT PROCESSING

The complete architecture

Software Architecture (H09B5a and H07Z9a) – Part 2b

Jeroen Reinenbergh (r0460600)
Jonas Schouterden (r0260385)

Academic year 2014–2015

Contents

1	Introduction	3
2	Overview	3
2.1	Architectural decisions	3
2.2	Discussion	5
3	Client-server view (UML Component diagram)	5
3.1	Main architectural decisions	7
3.1.1	Av1a & Av2a: Notifying the appropriate operator within 1 minute	7
3.1.2	Av1b: Storing the status of an individual job	8
3.1.3	Av2b: Temporarily storing documents that should be delivered via the personal document store for at least 3 hours in case the PDSDB is unavailable and providing a clear message to users in order to fail gracefully	8
3.1.4	Av3: Zoomit failure	8
3.1.5	P2: Document lookups	8
3.1.6	P3: Status overview for customer administrators	8
3.1.7	M1: New type of document: bank statements	8
3.1.8	M2: Multiple print & postal services	8
3.1.9	M3: Dynamic selection of the cheapest of print & postal services	8
4	Decomposition view (UML Component diagram)	8
4.1	DeliveryFunctionality	8
4.2	DocumentDB	8
4.3	DocumentGenerationManager	9
4.4	DocumentStorageFunctionality	9
4.5	JobManager	9
4.6	LinkMappingFunctionality	10
4.7	PDSDB	10
4.8	UserFunctionality	11
5	Deployment view (UML Deployment diagram)	11
6	Scenarios	12
6.1	Scenario 1	13
6.2	A registered recipient logs in	13
6.3	A registered recipient logs out	13
6.4	Updating a document template	14
6.5	Verifying a session	14
6.6	Initiating document processing	16
6.7	Consult document in personal document store	16
A	Element catalog	20
A.1	AuthenticationHandler	20
A.2	BillingManager	21
A.3	ChannelDispatcher	21
A.4	CustomerOrganizationClient	22
A.5	CustomerOrganizationFacade	22
A.6	Completer	23
A.7	DocumentDB	24
A.8	DocumentDBShard	24
A.9	DocumentDBShardingManager	24
A.10	DocumentGenerationManager	25
A.11	DocumentStorageCache	26
A.12	DocumentStorageFunctionality	26
A.13	DocumentStorageManager	27
A.14	DeliveryFunctionality	27
A.15	EDocsAdminClient	28

A.16 EmailChannel	28
A.17 EmailFacade	29
A.18 Generator	29
A.19 GeneratorManager	30
A.20 JobDBShard	30
A.21 JobDBShardingManager	30
A.22 JobFacade	31
A.23 JobManager	32
A.24 KeyCache	32
A.25 LinkMappingDB	32
A.26 LinkMappingManager	33
A.27 LinkMappingFunctionality	33
A.28 PDSDB	34
A.29 PDSDBReplica	35
A.30 PDSFacade	35
A.31 PDSLlongTermDocumentManager	36
A.32 PDSReplicationManager	37
A.33 Print&PostalServiceChannel	37
A.34 Print&PostalServiceFacade	38
A.35 NotificationHandler	38
A.36 OtherDB	38
A.37 RecipientClient	40
A.38 RawDataHandler	41
A.39 RecipientFacade	41
A.40 RegistrationManager	42
A.41 Scheduler	42
A.42 SessionDB	43
A.43 TemplateCache	44
A.44 UserFunctionality	44
A.45 ZoomitChannel	45
A.46 ZoomitFacade	45
B Defined data types	46

1 Introduction

The goal of this project was to design a document processing system. In this document, we describe the final architecture, which was designed based on the requirements from the domain analysis and the priorities of these requirements given by our stakeholders. The provided initial architecture was used as a starting point to achieve this. Section 2 lists the architectural decisions for all non-functional requirements and discusses the final architecture. Section 3 provides and discusses the main context and decomposition diagrams of our architecture (i.e. the context and primary diagram of the component-and-connector view) along with a discussion of the main architectural decisions involved. Section 4 provides and discusses the more fine-grained decompositions of some of the major components in the main decomposition. Section 5 provides and discusses the deployment of the components of the component-and-connector view on physical nodes. Finally, Section 6 illustrates how our architecture accomplishes the most important functionality and data flows using sequence diagrams. Afterwards, Appendix A lists and describes all components of the component-and-connector view and their interfaces and Appendix B lists and describes the data types used in these interfaces.

2 Overview

This section gives a high-level but complete overview of the system: it lists the design decisions for all non-functional requirements and provides a discussion concerning the strong and weak points of the architecture.

2.1 Architectural decisions

In this section, we give an overview of the architectural decisions made in our architecture in order to achieve the requirements given in the domain analysis and the residual drivers given in the provided initial architecture. However, we will not repeat any decisions that were already documented in this provided initial architecture as they can be found there.

Av1a & Av2a: Notifying the appropriate operator within 1 minute *Av1a* and *Av2a* require the system to notify the an eDocs operator within 1 minute. To achieve this, the `NotificationHandler` handles all incoming notifications and forwards all those that are destined for an eDocs operator to the `EDocsAdminFacade`, which in turn delivers these to the external `EDocsAdminFacade`. To ensure that notifications get sent to the eDocs operator in a timely fashion, both the `EDocsAdminFacade` and the `NotificationHandler` are deployed on the same node, which is separated from all other nodes to increase performance. For more details, we refer to Section 3.1.1 and Section 5.

Av1b: Storing the status of an individual job *Av1b* requires the system to store the status of an individual job. To achieve this, the `JobManager` accepts all read and write requests regarding the storage of job data, including status information. For more details, we refer to Section 3.1.2 and Section 4.5.

Av2b: Temporarily storing documents that should be delivered via the personal document store for at least 3 hours in case the PDSDB is unavailable and providing a clear message to users in order to fail gracefully *Av2b* requires (1) documents that should be delivered via the personal document store to be cached for at least 3 hours in case of unavailability and (2) providing a clear message to the recipient in this case. To achieve this, we introduced the `DocumentStorageCache`: this component temporarily stores the `DocumentIds` and corresponding `UserIds` of all generated documents that are to be delivered to the PDSDB during downtime of the latter component up to a maximum of 3 hours. When the PDSDB turns operational again, it notifies the `DocumentStorageManager` with a sign of life, which in turn retrieves all documents and their corresponding meta data from the `DocumentDB` using the previously mentioned ids and subsequently stores them in the PDSDB after conversion of the `DocumentMetaData`.

Note that the user therefore perceives a maximum total downtime of 3 hours and that the `RecipientFacade` is in charge of presenting the user with a clear error message after having performed a read request in his or her behalf that has failed due to the unavailability of the PDSDB. However, during the time needed for the `DocumentStorageManager` to transfer all documents and meta data that the cached ids refer to, it is possible that the user is not able to access (part of) his or her documents via the PDSDB, more specifically those that are still being transferred. Since the requirements do not demand the user to be able to regain full functionality of his or her personal document store at once, finding a better alternative for this gradual revival of the PDSDB

is out of scope.

Finally, we would like to stress that the `DocumentStorageManager` and the `PDSDB` are necessarily deployed on different nodes for the former to be able to do its caching job while the latter is not operational during downtime. More precisely, the `DocumentStorageManager` implicitly pings the `PDSDB` when storing document data in it. The echo message then, in turn, consists of the write confirmation that is subsequently received. If one of those writes should fail, all subsequent writes are internally converted into pairs of `DocumentId-UserId` writes to the aforementioned cache. Once the `PDSDB` is operational again and all cached ids are processed, subsequent writes to the `PDSDB` will no longer be redirected through the `DocumentStorageCache`.

For more details, we refer to Section 3.1.3, Section 4.4 and Section 5.

Av3: Zoomit failure *Av3* requires the system to be able to (1) autonomously detect when an invoice is not accepted by Zoomit, (2) temporarily store at least 2 days of these invoices locally until they are accepted by Zoomit, (3) keep on retrying to deliver a failed invoice to Zoomit in a proper fashion and (4) notify an eDocs operator after 5 failed attempts of at least 10 different invoices. To achieve this, the `ZoomitFacade` is able to receive a delivery confirmation from the `ZoomitChannel` (i.e. the invoice was accepted by Zoomit). Furthermore, the `ZoomitDeliveryCache` caches all `Documents` and corresponding `DeliveryInfo` for which the `ZoomitFacade` did not (yet) receive a delivery confirmation up to a maximum of 2 days of documents.

The `ZoomitFacade` then periodically retries to deliver the next document in the cached list of `Document-DeliveryInfo` pairs (note that the `ZoomitDeliveryCache` thus stores its entries in a circular manner) to the `ZoomitChannel` in a proper fashion by using exponential back-off. Once the `ZoomitFacade` has received a delivery confirmation regarding one of those cached documents, the corresponding `Document-DeliveryInfo` pair is deleted from the `ZoomitDeliveryCache` and the `ZoomitFacade` tries to deliver every other document in the list to the `ZoomitChannel` right after, disregarding the previously employed method of exponential back-off (because the external `ZoomitChannelNode` on which the `ZoomitChannel` resides, is now presumed to be operational again).

Finally, the `Document-DeliveryInfo` pairs in the `ZoomitDeliveryCache` need to be stored along with a counter that the `ZoomitFacade` increments after each failed attempt to send the respective document. Additionally, the `ZoomitFacade` keeps track of all cached pairs whose counters have reached the value of 5 and sends a notification to an eDocs operator through the `NotificationHandler`. For more details, we refer to Section 3.1.4, Section 3.1.1, Section 4.1 and Section 5.

P2: Document lookups *P2* requires the system to be able to (1) respond to all incoming document requests in a timely fashion, (2) throttle excessive requests when the arrival rate is larger than a certain request specific value (determined by the requirements) so that the non-excessive ones continue to be handled in a timely fashion and (3) make sure that the performance of all other functionality of the system remains unaffected in case of a large number of requests. To achieve this, all potential bottleneck components that support the document lookup process are deployed on separate nodes. More precisely, the `RecipientFacade` and the `LinkMappingFunctionality` both reside on the same exclusive node as to increase the performance of these components and the communication between them (i.e. when a recipient has presented the former component with a link for which it has to determine the document it maps to). The `PDSFacade` is deployed on the `PDSDBManagerNode` (discussed in the provided initial architecture), along with the `PDSLongTermManager` and the two `PDSReplicationManagers`, as to increase communication performance between this `PDSFacade` and the `PDSDB` supercomponent. Notice that *Av1* already demands the `PDSDB` to be deployed separate from all other functionality according to the provided initial architecture and that *P2* can rely on this design decision to deliver an optimally increased (taking the availability constraints of *Av1* into account) performance of this supercomponent.

Because the `DocumentDB`, like the `PDSDB`, stores a large amount of documents and correspondingly handles a large amount of document requests, its documents are partitioned across several `DocumentDBShards` that each reside on their own `DocumentDBShardNode`. This sharding technique results in a performance increase that is statistically equivalent to the one that is achieved by active replication, but it is significantly less costly regarding storage costs (which is more of a concern for documents than for other kinds of data). In order to further increase performance of the `DocumentDB`, the `DocumentDBShardingManager`, which is in charge of managing all `DocumentDBShards` and forwarding read and write requests to the appropriate ones, is also deployed on an individual node.

Finally, both the `DocumentDB` and the `PDSDB` have the ability to throttle excessive read requests from the `RecipientFacade` (all other requests remain unaffected) when the rate of these incoming requests exceeds a certain request specific value (determined by the requirements). Also note that the performance of all other functionality of the system remains unaffected even at this maximum, since all potential bottleneck

components that support the document lookup process are deployed on separate nodes, effectively increasing their performance (as discussed above) and thereby providing both an increased level of performance for the document lookup process and a sufficient level of performance for all other functionality. For more details, we refer to Section 3.1.5, Section 3, Section 4.8, Section 4.7, Section 4.2 and Section 5.

P3: Status overview for customer administrators *P3* requires (1) the status overview of all document processing jobs initiated by a certain customer organization to be delivered to the respective customer administrator(s) in a timely fashion (i.e. the provided status of a document processing job should be consistent up to 1 minute ago) and (2) the construction of this status overview not to hinder other parts of the system. To achieve this, all potential bottleneck components that support the status overview process are deployed on separate nodes as to increase their performance. More precisely, the `CustomerOrganizationFacade` is deployed on its own node (together with the `SessionDB` and the `AuthenticationHandler` for reasons discussed elsewhere) and the `JobManager` supercomponent is deployed separate from all other components in the system too. The latter is responsible for storing all jobs, which are partitioned across several `JobDBShards` that each reside on their own `JobDBShardNode`. This sharding technique results in a performance increase that is statistically equivalent to the one that is achieved by active replication, but it is less costly regarding storage costs. In order to further increase performance of the `JobManager` and the corresponding speed at which job statuses can be requested, the `JobDBShardingManager`, which is in charge of managing all `JobDBShards` and forwarding read and write requests to the appropriate ones, is also deployed on an individual node.

Note that the performance of all other functionality of the system remains unaffected during the construction of such a status overview, since all potential bottleneck components that support this status overview process are deployed on separate nodes, effectively increasing their performance and thereby providing both an increased level of performance for the status overview process and a sufficient level of performance for all other functionality.

Also note that the `JobDBShardManager` and corresponding `JobDBShards` do not distinguish between age of jobs

For more details, we refer to Section 3.1.5, Section 3, Section 4.8, Section 4.7, Section 4.2 and Section 5.

M1: New type of document: bank statements

M2: Multiple print & postal services

M3: Dynamic selection of the cheapest of print & postal services

2.2 Discussion

GEEN JOBFACADE MEER Use this section to discuss your architecture in retrospect. For example, what are the strong points of your architecture? What are the weak points? Is there anything you would have done otherwise with your current experience? Are there any remarks about the architecture that you would give to your customers? Etc.

Alternative for DocumentStorageCache We could just as well do without the previously introduced `DocumentStorageCache` by letting the `DocumentStorageManager` actively look for documents in the `DocumentDB` that are not yet, but should be, present in the `PDSDB` upon revival of the `PDSDB`. A clear advantage of this approach is that the downtime of the `PDSDB` component is no longer restricted by the aforementioned 3-hour cache. An important disadvantage, however, lies in the fact that the `DocumentStorageManager` is burdened with a significant amount of extra work and will require more expensive hardware to cope with this. Since the support for a longer downtime of the `PDSDB` is out of scope, this approach was not chosen.

3 Client-server view (UML Component diagram)

Context diagram The context diagram of the component-and-connector view is given in Figure 1. As shown, six distinct types of external components communicate with the system: (i) the `EDocsAdminClient` represents a client device of an administrator of eDocs that communicates with the eDocs System, (ii) the `CustomerOrganizationClient` represents a client device of a Customer Organization (i.e. Customer Administrator and Customer Information System) that communicates with the eDocs System, (iii) the `RecipientClient` represents a client device of an unregistered or registered recipient of eDocs that communicates with the

eDocs System, (iv) the ZoomitChannel represents the servers of Zoomit to which documents can be sent for further delivery, (v) the Print&PostalServiceChannel represents the servers of a print & postal service to which documents can be sent for further delivery, (vi) the EmailChannel represents a mail server of an e-mail provider to which documents can be sent for further delivery. Notice that each external interface on the client side is provided by a separate facade component and that each delivery channel interfaces with the DeliveryFunctionality component.

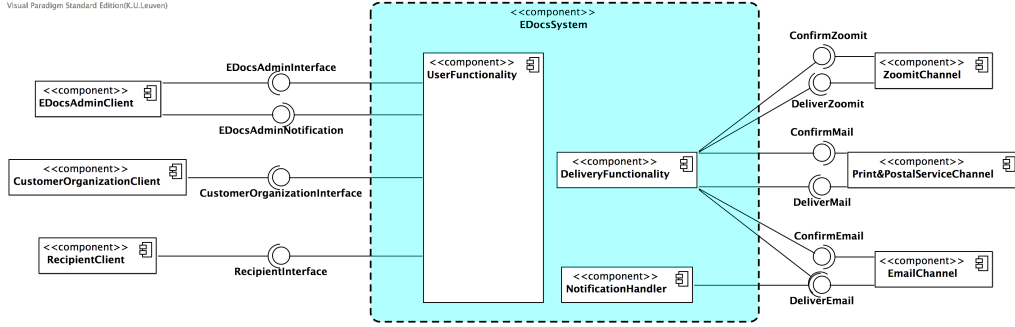


Figure 1: Context diagram for the client-server view.

Primary diagram The main component-and-connector view is shown in Figure 2. All components that were decomposed further are highlighted. The DeliveryFunctionality is further decomposed in Section 4.1, the DocumentDB in Section 4.2, the DocumentGenerationManager in Section 4.3, the DocumentStorageFunctionality in Section 4.4, the JobManager in Section 4.5, the LinkMappingFunctionality in Section 4.6, the PDSDB in Section 4.7 and the UserFunctionality in Section 4.8. The descriptions of all components in Figure 2, their sub-components and their interfaces are given in Appendix A.

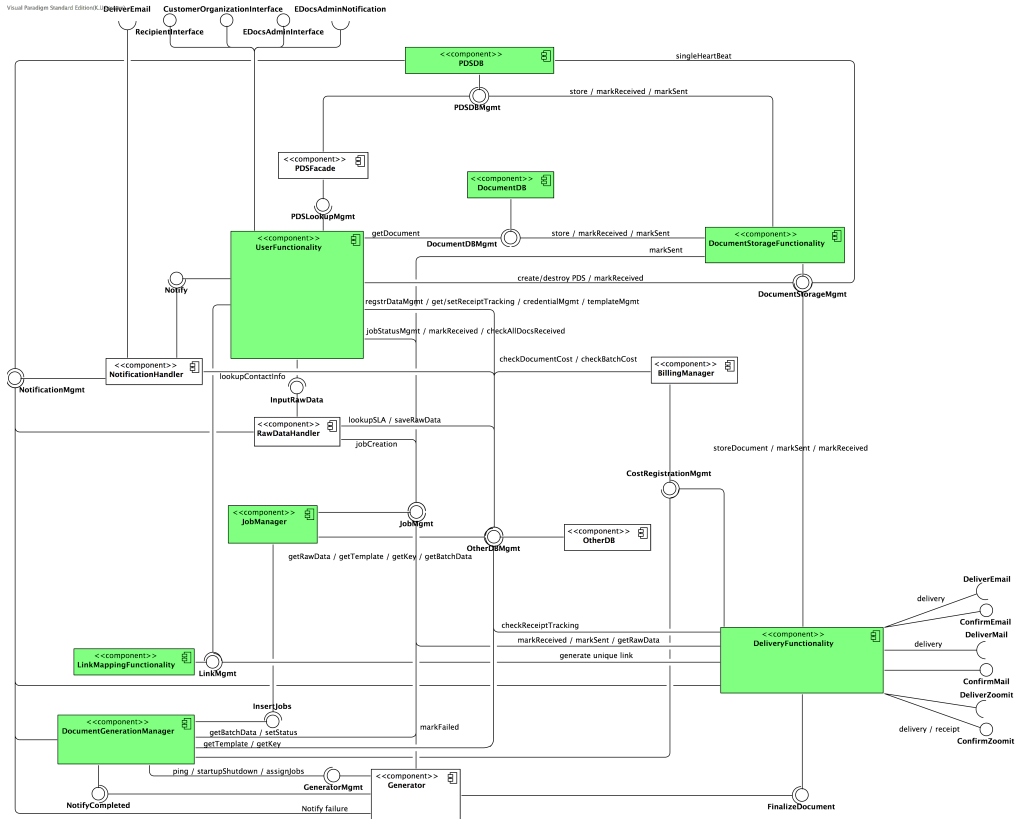


Figure 2: Primary diagram of the client-server view.

3.1 Main architectural decisions

Discuss your architectural decisions for the most important requirements in more detail using the components of the client-server view. Pay attention to the solutions that you employed and the alternatives that you considered. The explanation here must be self-contained and complete. Imagine you had to describe how the architecture supports the core functionality to someone that is looking at the client-server view only. Hide unnecessary details (these should be shown in the decomposition view).

UserFacade We introduce this component to be able to distinguish between registered recipients and unregistered recipients in the first steps of the document lookup process. The **UserFacade** also maps incoming links to documents that are located either in the **PDSDB** or in the **DocumentDB** by first fetching the correct mapping from the **LinkMappingDB** mentioned above. Another responsibility of the **UserFacade** is marking documents as received, since this is the last internal component through which the document is passed before the requesting recipient actually receives it and failure of another component in the document lookup pipeline can no longer prevent this from happening.

PDSFacade This component handles all read requests that are intended for the personal document store. This extra level of indirection calculates and aggregates all intermediate query results as to relieve the **PDSDB** of this burden, which is also the reason why both components are best deployed on different nodes. Upon querying the personal document store, this component first collects the recipient's documents' meta data and subsequently performs queries on this collection. This approach introduces no significant overhead, because only those documents that are needed, will be fetched in their entirety. It is also the responsibility of the **PDSFacade** to check the recipient's ID against the recipient ID that can be retrieved from the document meta data and to only pass on those documents for which they are a match.

DocumentStorageManager The storage of documents is handled by the **DocumentStorageManager**, which receives generated documents from the intermediary **OtherFunctionality2** component and stores them in either the **DocumentDB** (for unregistered recipients) or in both the **DocumentDB** and the **PDSDB** (for registered recipients) according to the method that is invoked on it. The necessity for this component follows from the fact that synchronisation is needed between the two storage components mentioned above. Note that the **PDSDocMgmt** interface in the **PDSDB** component now requires an extra method to save documents.

LinkManager Motivatie voor LinkManager: Checks expiration date ALS DAT NODIG IS – reason: links naar de pdsdb vervallen niet (zolang de gebruik geregistreerd is) LinkManager maps link to (document ID, place where the document is stored)-pairs – REASON: the unique link has two possible sources: an e-mail to an unregistered recipient or an email to a registered recipient. For an unregistered recipient, the RecipientFacade must look with the documentid for the document in the documentDB. For a registered recipient, the RecipientFacade must look with the documentid for the document in the PDSDB. (Mogelijk een boolean ofzo) Does NOT do mapping removal after x years – there has to be a notification when the link has expired

3.1.1 Av1a & Av2a: Notifying the appropriate operator within 1 minute

Describe the design choices related to *ReqX* together with the rationale of why these choices were made.

Alternatives considered

Alternative(s) for choice 1 Explain what alternative(s) you considered for this design choice and why they were not selected.

3.1.2 Av1b: Storing the status of an individual job

3.1.3 Av2b: Temporarily storing documents that should be delivered via the personal document store for at least 3 hours in case the PDSDB is unavailable and providing a clear message to users in order to fail gracefully

3.1.4 Av3: Zoomit failure

3.1.5 P2: Document lookups

Sharding for DocumentDB for P2 Note that there must also be a (sub)component monitoring all requests to the shards, in order to throttle excessive requests when necessary. Since this is a rather simple task and the (sub)component must be aware of the details concerning the sharding architecture to efficiently fulfil its purpose, this functionality is delegated to the **ShardingManager** itself. Finally, this **ShardingManager** is also responsible for implicitly pinging all **DocumentDBShards** upon reading and writing in them.

3.1.6 P3: Status overview for customer administrators

3.1.7 M1: New type of document: bank statements

3.1.8 M2: Multiple print & postal services

3.1.9 M3: Dynamic selection of the cheapest of print & postal services

4 Decomposition view (UML Component diagram)

Discuss the decompositions of the components of the client-server view which you have further decomposed.

4.1 DeliveryFunctionality

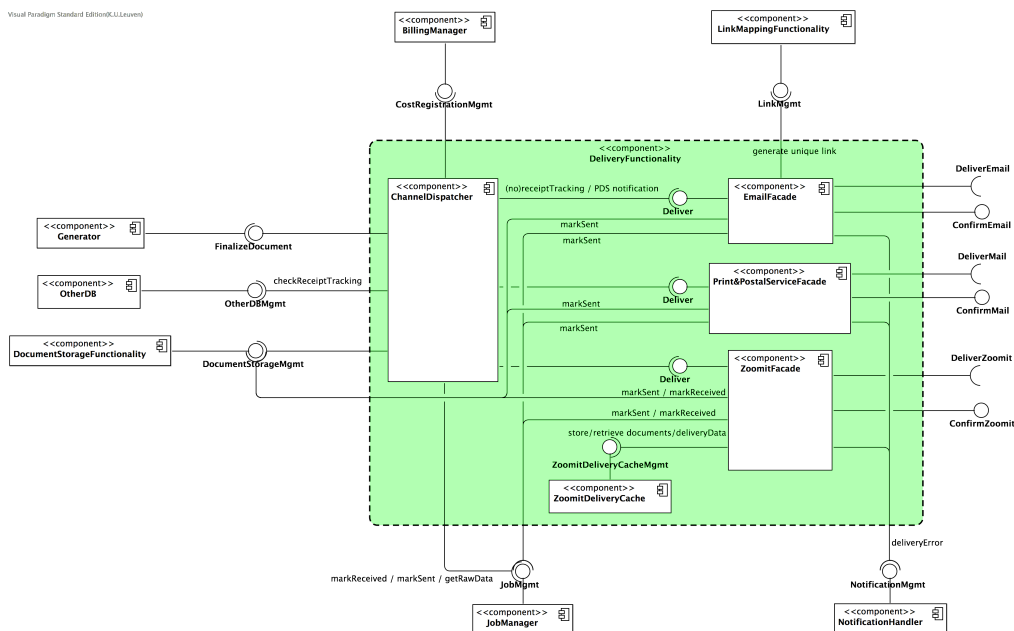


Figure 3: Decomposition of DeliveryFunctionality

Describe the decomposition of ComponentX and how this relates to the requirements.

4.2 DocumentDB

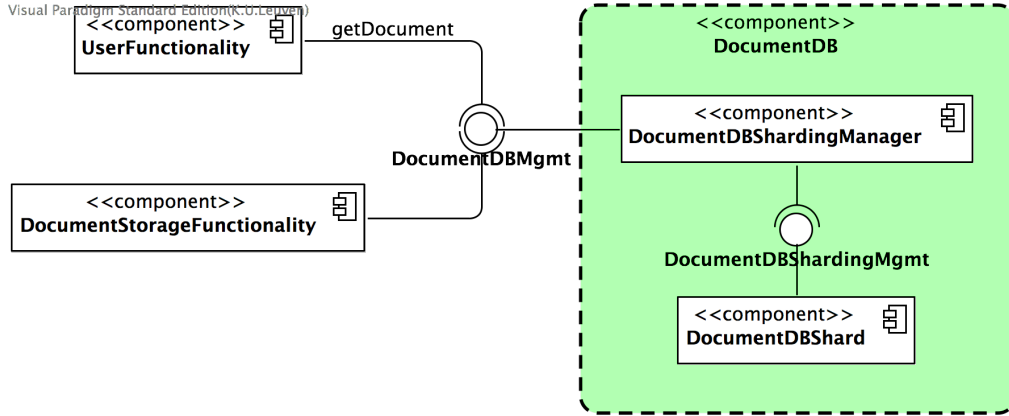


Figure 4: Decomposition of DocumentDB

4.3 DocumentGenerationManager




Figure 5: Decomposition of DocumentGenerationManager

4.4 DocumentStorageFunctionality

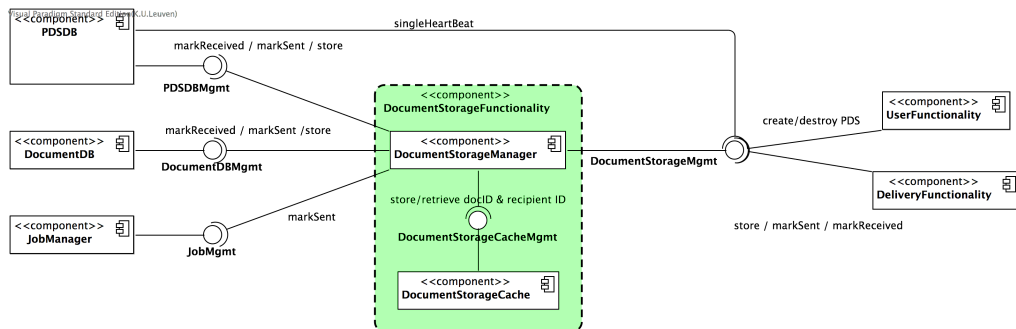


Figure 6: Decomposition of DocumentStorageFunctionality

4.5 JobManager

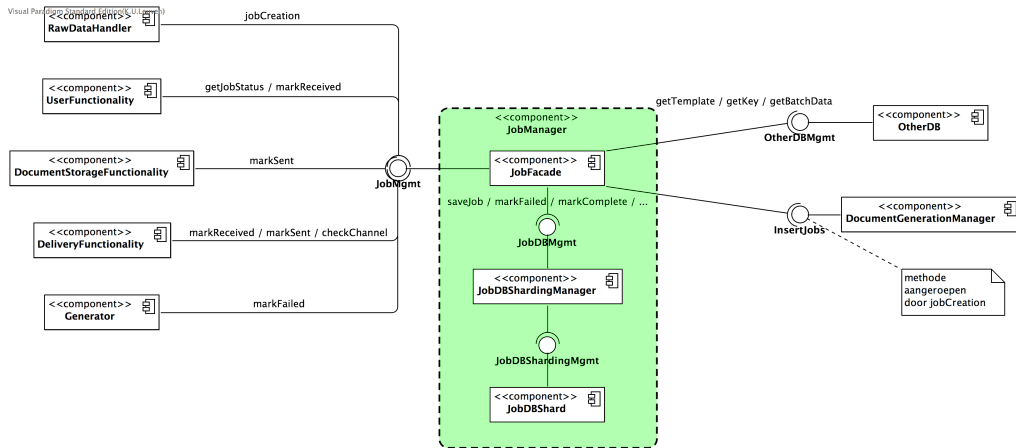


Figure 7: Decomposition of JobManager

4.6 LinkMappingFunctionality

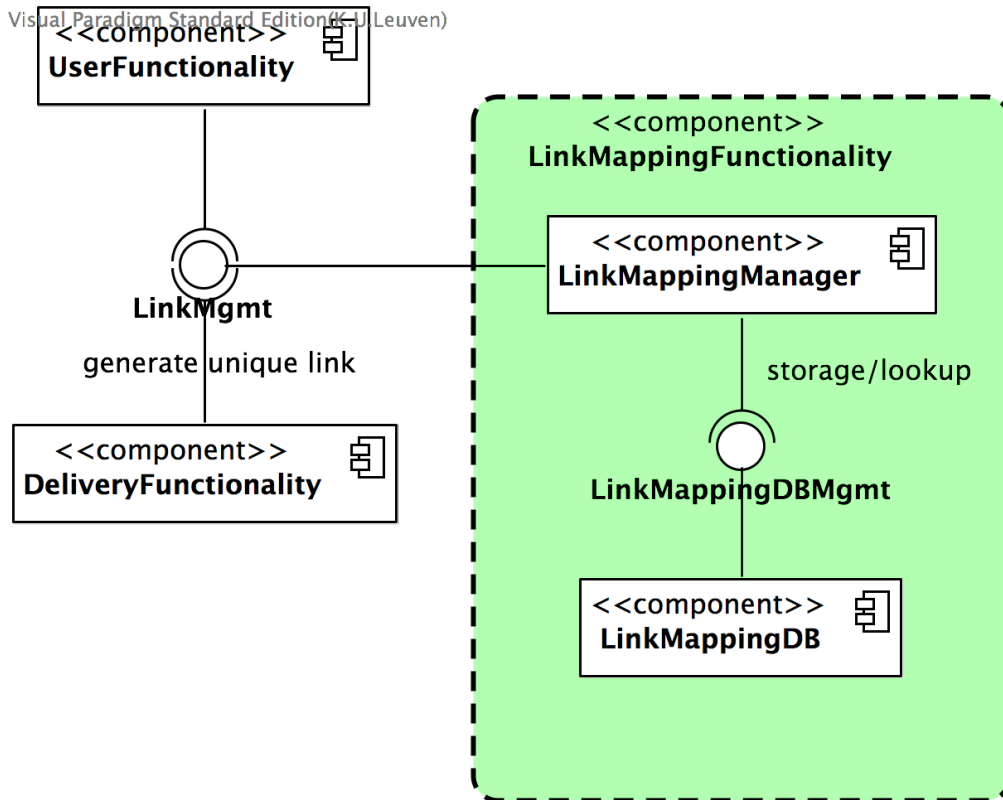


Figure 8: Decomposition of LinkMappingFunctionality

4.7 PDSDB




Figure 9: Decomposition of PDSDB

4.8 UserFunctionality

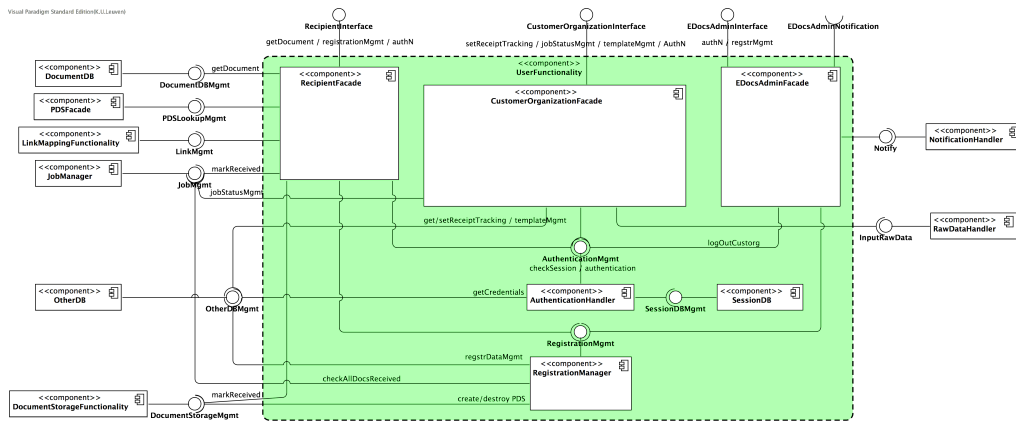


Figure 10: Decomposition of UserFunctionality

5 Deployment view (UML Deployment diagram)

Describe the context diagram for the deployment view. For example, which protocols are used for communication with external systems and why?

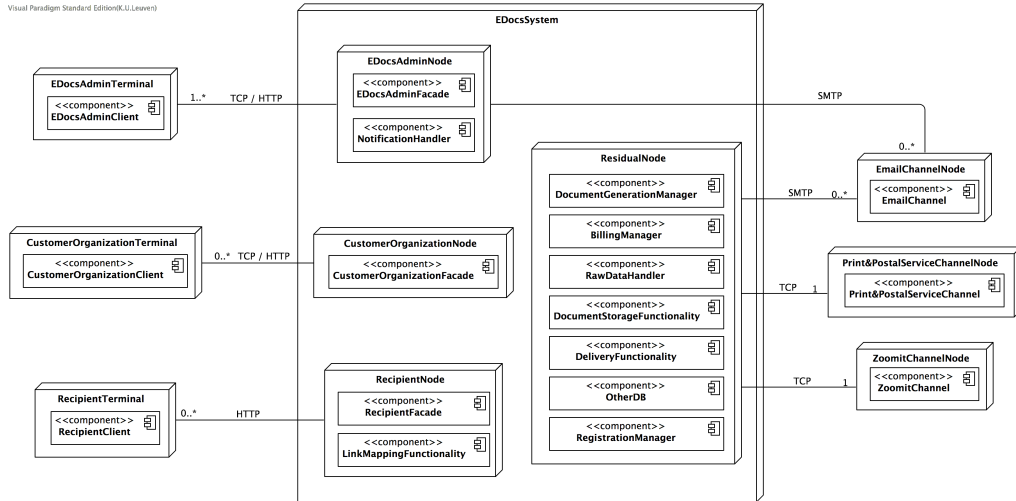


Figure 11: Context diagram for the deployment view.

The primary deployment diagram itself and accompanying explanation. Pay attention to the parts of the deployment diagram which are crucial for achieving certain non-functional requirements. Also discuss any alternative deployments that you considered.

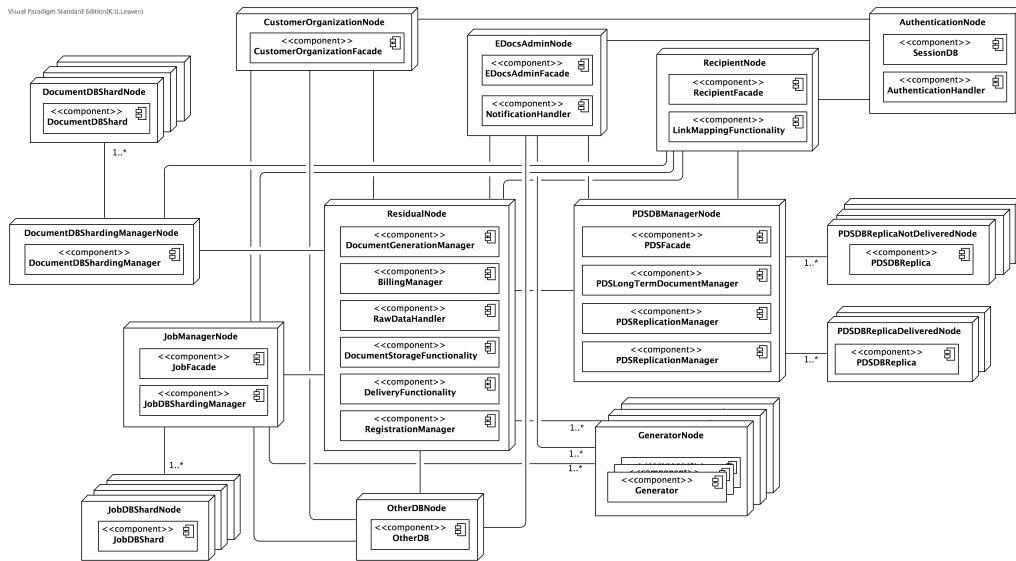


Figure 12: Primary diagram for the deployment view.

6 Scenarios

Illustrate how your architecture fulfills the most important data flows. As a rule of thumb, focus on the scenario of the domain description. Describe the scenario in terms of architectural components using UML Sequence diagrams and further explain the most important interactions in text. Illustrating the scenarios serves as a quick validation of the completeness of your architecture. If you notice at this point that for some reason, certain functionality or qualities are not addressed sufficiently in your architecture, it suffices to document this, together with a rationale of why this is the case according to you. You do not have to further refine your architecture at this point.

6.1 Scenario 1

Shortly describe the scenario shown in this subsection. Show the complete scenario using one or more sequence diagrams.




Figure 13: The system behavior for the first scenario.

6.2 A registered recipient logs in

Figure 14 depicts the sequence diagram of a Registered Recipient logging in, according to *UC1: Log in*. The Registered Recipient provides his details to the **RecipientFacade**. The **Authentication** compares these credentials to those stored in the **OtherDB**. If the given credentials match the stored credentials, a new session is opened in the **SessionDB** and the corresponding session identifier is returned to the registered recipient. Otherwise, an **InvalidCredentialsException** is thrown.

Note that the use case *UC1: Log in* also has the Customer Organization as a primary actor. The login procedure for customer organizations is almost identical to the procedure for registered recipient, which is why we do not provide a separate sequence diagram for this primary actor. Compared to figure 14, the Customer Organization sends its login request to the **CustomerOrganizationFacade** instead of the **RecipientFacade**. Also, another method call to **OtherDB** is used to ask for the credentials of a customer organization.




Figure 14: The login behaviour of a Registered Recipient.

6.3 A registered recipient logs out

Figure 15 shows the behaviour of a Registered Recipient logging out, according to *UC2: Log out*. The Registered Recipient provides his recipient id and session id to the **RecipientFacade**. The **AuthenticationHandler**

first verifies the session (figure 15). If the session is valid, it logs out the Registered Recipient. Note that the logout procedure for Customer Organizations is almost the same. The only difference is that the Customer Organization sends his logout request to the **CustomerOrganizationFacade** instead of the **RecipientFacade**. Because of this reason, we do not give a separate sequence diagram for the Customer Organization logging out.




Figure 15: The logout behaviour of the Registered Recipient.

6.4 Updating a document template

Figure 16 shows the flow when a Customer Administrator updates the template used for generating documents of a given type. It corresponds to *UC20: Update document template*. It consists of two main parts. First, the customer administrator asks what the possible document types are that the customer organization is allowed to generate. The system verifies whether the session of the customer administrator is valid (as detailed in figure 17 and then returns a list of the allowed document types using the **OtherDB**. Together with the allowed document types, it returns the time when the currently stored templates corresponding to the document types were uploaded.

Secondly, the customer administrator provides the document type it wants to update and the new template to the **CustomerOrganizationFacade**. After verifying the validity of the session again (figure 17), the **CustomerOrganization** asks the **OtherDB** to check whether the document type provided by the customer organization is valid and allowed. If it is not allowed or invalid, an **InvalidDocumentTypeException** gets thrown. Otherwise, the **CustomerOrganizationFacade** stores the template in the **OtherDB** together with the time when it has received the template, the document type and the customer organization id.

6.5 Verifying a session

Figure 17 how it is verified whether a session is valid. The **AuthenticationHandler** uses the **SessionDB** to verify whether each incoming session identifier belongs to an existing session. If it belongs to an existing session, the corresponding session attributes are returned to the caller, otherwise an **NoSuchSessionException** is thrown.




Figure 16: The system behavior for updating a document template.




Figure 17: verifySession: sequence diagram depicting the verification whether a session is valid.

6.6 Initiating document processing

Figure 18 shows how a customer organization initiates a batch of document processing jobs. Recurring and non-recurring batches are indicated using the `isRecurring` boolean value. If the batch is non-recurring, the customer organization can indicate another priority than the default priority of that customer organization if the batch could be handled with another priority than the default priority. For recurring batches, it does not matter if a priority is given, as the default priority will be used. the customer organization also delivers the raw data as a `rawDataPackage` and the meta-data of the batch to the `CustomerOrganizationFacade`. After verification of the session identifier (figure 17), the `CustomerOrganizationFacade` will put a `TimeStamp` of the time when this method is called in the `batchMetaData` and forward all the information to `RawDataHandler`, while also indicating if the batch is (non-)recurring.

Figure 19 depicts the first part of the method call in `RawDataHandler`, which will verify whether the raw data is valid. First, `RawDataHandler` checks whether the customer organization is allowed to generate the given document type. If it is allowed, the received raw data packet will be validated, e.g. there will be checked whether the content of the received Excel file can be read or whether a received XML file is correctly formatted. If it is not valid, an exception is thrown. Otherwise, depending on if the batch is recurring, the `RawDataHandler` will query the SLA in the `OtherDB` to check if the number of raw data entries in the raw data package is allowed. If the number of entries exceeds this maximum, an exception is thrown. Next, each individual entry will be checked on validity, so the `RawDataHandler` knows which `RawData` entries are valid and which are not.

After verification, figure 18 shows that the `RawDataHandler` will only use the given priority if the batch is non-recurring. Otherwise, it will use the default priority. The deadline of the document processing jobs is calculated and put into the `BatchMetaData`, after which the `RawData` entries and `BatchMetaData` are stored in the `OtherDB`. The `RawDataHandler` asks the `JobFacade` to create jobs for the valid raw data entries and to initiate document processing for those jobs.

Figure 20 shows the `JobFacade` creating and storing jobs for the valid raw data entries, after which it inserts these jobs into the `Scheduler`.

6.7 Consult document in personal document store

Shortly describe the scenario shown in this subsection. Show the complete scenario using one or more sequence diagrams.

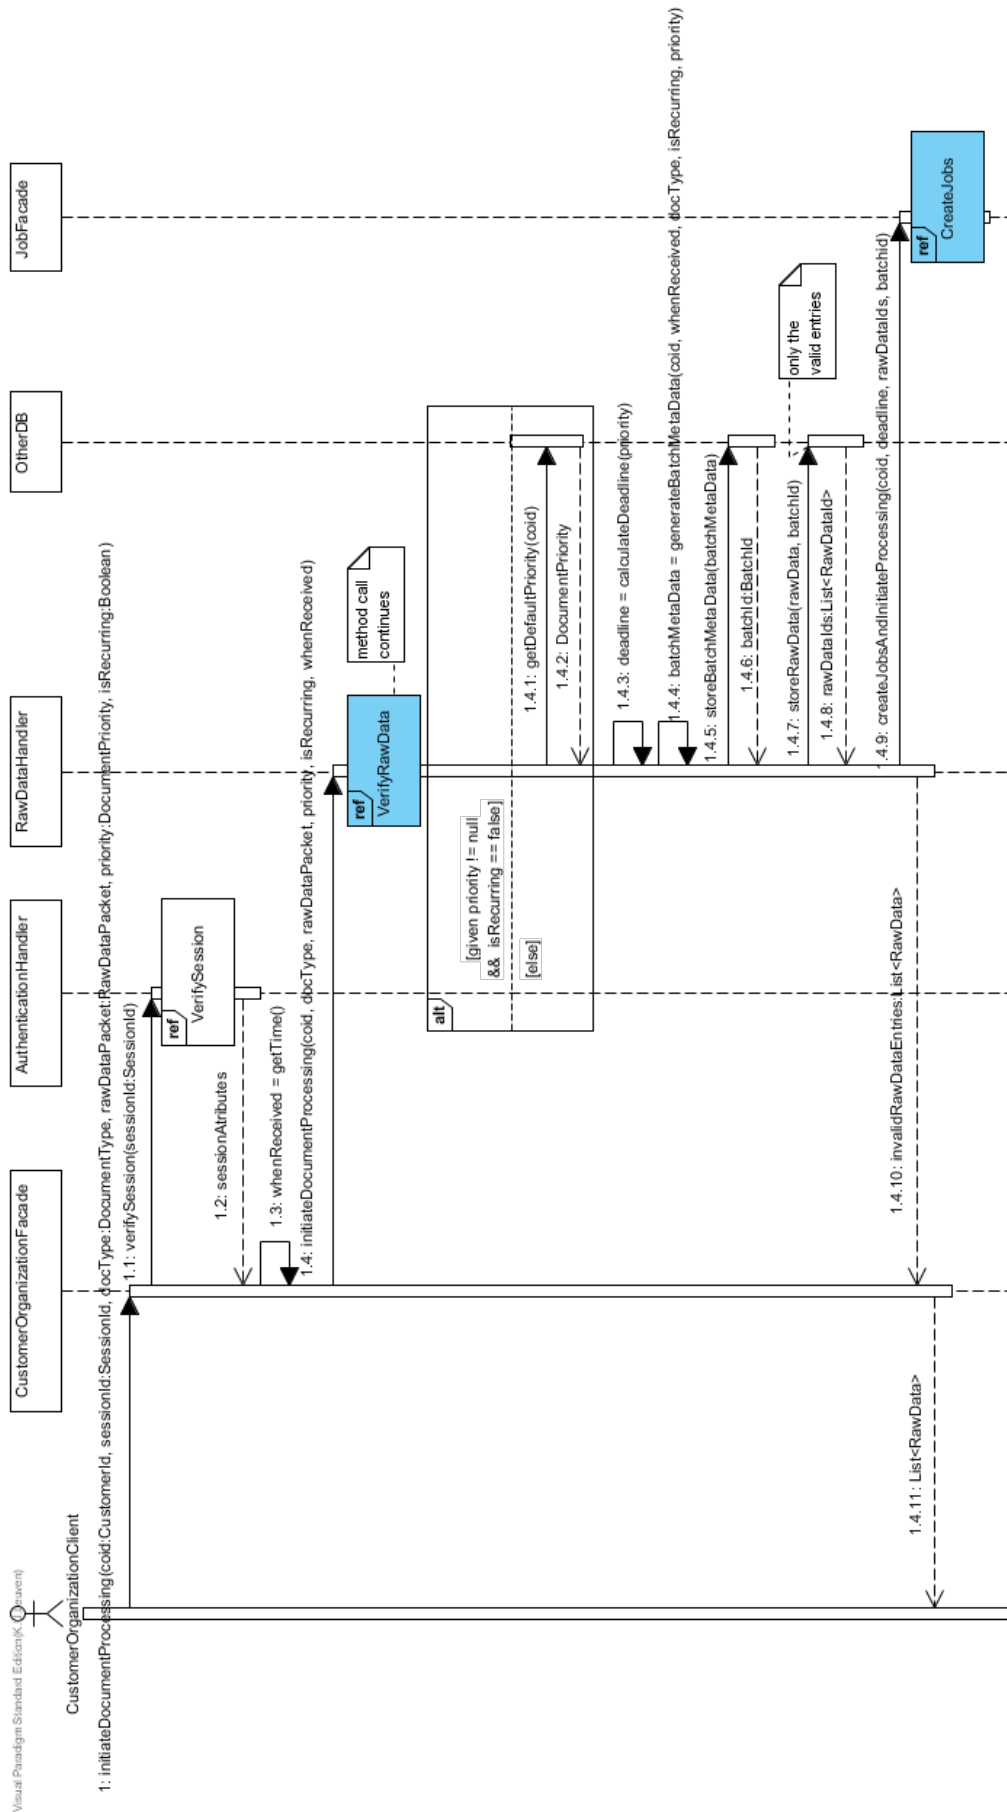


Figure 18: VerifyRawData: verification whether the raw data is valid.




Figure 19: VerifyRawData: verification whether the raw data is valid.




Figure 20: CreateJobs: the JobManager creates jobs and forwards them to be scheduled by the Scheduler.




Figure 21: A Registered Recipient consults a document in his or her personal document store.

A Element catalog

In this section, we list all the components and the interfaces they provide. Per component, we describe its responsibilities, declare its super-component (if any) and list its sub-components (if any).

List all components and describe their responsibilities and provided interfaces. Per interface, list all methods using a Java-like syntax and describe their effect and exceptions if any. List all elements and interfaces alphabetically for ease of navigation.

A.1 AuthenticationHandler

- **Description:** The `AuthenticationHandler` is responsible for authenticating Registered recipients and Customer organizations. The architecture does not specify the means of authentication (e.g. the type of credentials). The credentials are stored in the `UserDB`.
- **Super-component:** `UserFunctionality`
- **Sub-components:** None

Provided interfaces

- `AuthN`
 - `RecipientId getId(SessionId sessionId) throws NoSuchSessionException`
 - * Effect: The `AuthenticationHandler` fetches and returns the Registered Recipient's identifier corresponding to the `sessionId` from the `SessionDB`.
 - * Exceptions:
 - `NoSuchSessionException`: Thrown if no session exists with the given identifiers, or if the session belongs to a customer organization.
 - `CustomerId getId(SessionId sessionId) throws NoSuchSessionException`
 - * Effect: The `AuthenticationHandler` fetches and returns the Customer Organization's identifier corresponding to the `sessionId` from the `SessionDB`.
 - * Exceptions:
 - `NoSuchSessionException`: Thrown if no session exists with the given identifier, or if the session belongs to a registered recipient.
 - `Boolean logout(SessionId sessionId)`
 - * Effect: The `AuthenticationHandler` will remove the session with the given id from the `SessionDB`. If no such session exists, nothing is changed and no exception is thrown.
 - * Exceptions: None
 - `SessionId login(Credentials credentials) throws InvalidCredentialsException`
 - * Effect: The `AuthenticationHandler` verifies the `credentials` using the `UserDB`. If they are correct, the `AuthenticationHandler` creates a new session using the `SessionDB`, stores the id of the user (i.e. the Registered Recipient id or Customer Organization id) as an attribute in this session and returns the id of the new session. The id of the user is present in the given credentials.
 - * Exceptions:
 - `InvalidCredentialsException`: Thrown if the given credentials are invalid.
- `CheckSession`
 - `Map<SessionAttributeKey, SessionAttributeValue> verifySession(SessionId sessionId) throws NoSuchSessionException`
 - * Effect: The `AuthenticationHandler` verifies whether a session with the given id exists in the `SessionDB` and if so, returns all its associated attributes.
 - * Exceptions:
 - `NoSuchSessionException`: Thrown if no session exists with the given identifiers.

A.2 BillingManager

- **Description:** The **BillingManager** is responsible for all billing tasks. This includes billing the Customer Organization for the generation and delivery of non-recurring document processing jobs.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- **CostRegistrationMgmt**
 - `void addDocumentGenerationCost(CustomerId cuId, JobId jobId, DocumentPriority priority)`
 - * Effect: The **BillingManager** adds the cost for generation documents corresponding to the job identified by `jobId` to the bill of the customer organization identified by `cuId`. The document had a priority `priority`. This method is called by the **GenerationManager**.
 - * Exceptions: None
 - `void addDocumentDeliveryCostByEmail(CustomerId cuId, JobId jobId)`
 - * Effect: The **BillingManager** adds the cost for delivering the document by e-mail to the bill of the customer organization identified by `cuId`. The document corresponds to the job identified by `JobId`. This method is called by the **ChannelDispatcher**.
 - * Exceptions: None
 - `void addDocumentDeliveryCostbyPrintAndPostalService(CustomerId cuId, JobId jobId)`
 - * Effect: The **BillingManager** adds the cost for delivering the document by e-mail to the bill of the customer organization identified by `cuId`. The document corresponds to the job identified by `JobId`. This method is called by the **ChannelDispatcher**.
 - * Exceptions: None
 - `void addDocumentDeliveryCostbyZoomit(CustomerId cuId, JobId jobId)`
 - * Effect: The **BillingManager** adds the cost for delivering the document by Zoomit to the bill of the customer organization identified by `cuId`. The document corresponds to the job identified by `JobId`. This method is called by the **ChannelDispatcher**.
 - * Exceptions: None

A.3 ChannelDispatcher

- **Description:** The **ChannelDispatcher** is responsible for choosing the correct delivery channel for a generated document. It also forwards the document to the **DocumentStorageManager**, which will store the document.
- **Super-component:** **Deliveryfunctionality**
- **Sub-components:** None

Provided interfaces

- **FinalizeDocument**

Note that the methods in this interface are made idempotent. The methods of this interface are called by **Generator** instances.

 - `void storeAndDeliverDocument(JobId jobid, Document doc)`
 - * Effect: The **ChannelDispatcher** will store the given document `document` and deliver it. This method is made idempotent. To filter duplicate method calls, it has the `JobId` of the document as an argument. This idempotence is to account for the case when a **Generator** fails after forwarding the document and before reporting completion to the **DocumentGenerationManager**. In this case, it can be that the **DocumentGenerationManager** restarts jobs for which a document has already been stored or delivered.
 - * Exceptions: None

- void generationError(JobId jobId, Error error)
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.4 CustomerOrganizationClient

- **Description:** The `CustomerOrganizationClient` is external to the eDocs system and represents a client device of a Customer Organization (i.e. Customer Administrator and Customer Information System) that communicates with the eDocs System.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - returnType1 operation1(ParamType param) throws SomeException
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - void operation2(ParamType2 param)
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.5 CustomerOrganizationFacade

- **Description:** The `CustomerOrganizationFacade` provides the main interface of the system to the Customer Organization (i.e. Customer Administrator and Customer Information System).
- **Super-component:** UserFunctionality
- **Sub-components:** None

Provided interfaces

- AuthN
 - Boolean logout(SessionId sessionId)
 - * Effect: The `AuthenticationHandler` will remove the session with the given id from the `SessionDB`. If no such session exists, nothing is changed and no exception is thrown.
 - * Exceptions: None
 - SessionId login(Credentials credentials) throws InvalidCredentialsException
 - * Effect: The `AuthenticationHandler` verifies the `credentials` using the `UserDB`. If they are correct, the `AuthenticationHandler` creates a new session using the `SessionDB`, stores the id of the user (i.e. the Registered Recipient id or Customer Organization id) as an attribute in this session and returns the id of the new session. The id of the user is present in the given credentials.
 - * Exceptions:
 - InvalidCredentialsException: Thrown if the given credentials are invalid.
- TemplateMgmt
 - List<Tuple<DocumentType,TimeStamp>> getPossibleDocumentTypes(SessionId sessionId, CustomerId cuId) throws NotAuthenticatedException

- * Effect: The `CustomerOrganizationClient` will return a list of the document types that the customer organization is allowed to generate. It also indicates the current template for each document type by returning for each document type the time when the last template was uploaded.
- * Exceptions:
 - `NotAuthenticatedException`: Thrown if the given session identifier is invalid.
- `Boolean updateDocumentTemplate(SessionId sessionId, CustomerId cuId, DocumentType documentType, Template template)` throws `NotAuthenticatedException`, `InvalidDocumentTypeException`
 - * Effect: The `CustomerOrganizationClient` will update the current template for the given `documentType` to the given `template` for the customer organization identified by `cuId`. Returns `true` when it succeeds.
 - * Exceptions:
 - `NotAuthenticatedException`: Thrown if the given session identifier is invalid.
 - `InvalidDocumentTypeException`: Thrown if the given document type is invalid or not allowed for the given customer organization.
- **InitiateDocumentProcessing**
 - `List<RawData> initiateDocumentProcessing(CustomerId cuid, SessionId sessionId, DocumentType docType, RawDataPacket rawDataPacket, DocumentPriority priority, Boolean isRecurring)` throws `NotAuthenticatedException`, `InvalidDocumentTypeException`, `InvalidRawDataPackageException`, `ToManyRawDataEntriesException`
 - * Effect: The `CustomerOrganizationFacade` gets an indication from a customer organization with customer identifier `cuid` to start a batch of document processing jobs. The documents to be generated should be of document type `docType`. The boolean `isRecurring` indicates whether the batch is recurring or non-recurring. The information necessary to generate the documents is given in a `RawDataPacket` containing all the info about the raw data entries of the batch. After verification of the session identifier, the `CustomerOrganizationFacade` generates a `TimeStamp` of the time when this method is called and forwards it together with most of its arguments to the `RawDataHandler`.
 - * Exceptions:
 - `NotAuthenticatedException`: Thrown if the given session identifier is invalid.
 - `InvalidDocumentTypeException`: Thrown if the given document type is invalid or not allowed for the given customer organization.
 - `InvalidRawDataPackageException`: Thrown if the raw data package contains an error (e.g. an invalid Excel file or incorrectly formatted XML).
 - `ToManyRawDataEntriesException`: Thrown if the batch is a recurring batch and contains more raw data entries than maximally allowed for this batch.

A.6 Completer

- **Description:** The `Completer` is responsible for fetching the raw data and applicable meta-data for a group of `JobIds` when a `Generator` instance requires a new group of jobs.
- **Super-component:** `DocumentGenerationManager`
- **Sub-components:** None

Provided interfaces

- **Complete**
 - `CompletePartialBatchData getComplete(BatchId batchId, List<JobId> jobIds)`
 - * Effect: The `Completer` fetches data needed by a `Generator` for generation of the documents corresponding to the `JobIds` belonging to the same batch, which is identified by `BatchId`.
 - * Exceptions: None

A.7 DocumentDB

- **Description:** The DocumentDB is responsible for actually storing all the documents. It stores documents regardless of the fact a document is also stored in the PDSDB. It receives read and write requests from the DocumentStoragManager.
- **Super-component:** None
- **Sub-components:** DocumentDBShardingManager and DocumentDBShard

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.8 DocumentDBShard

- **Description:** A DocumentDBShard is responsible for storing a partition of all the documents.
- **Super-component:** DocumentDB
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.9 DocumentDBShardingManager

- **Description:** The DocumentDBShardingManager manages the storage of the documents over multiple DocumentDBShards.
- **Super-component:** The DocumentDB
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.10 DocumentGenerationManager

- **Description:** The `DocumentGenerationManager` monitors the availability of the `Generator` components using the Ping interface. The `DocumentGenerationManager` keeps track of the jobs assigned to and being processed by the `Generators`. To minimize the overhead of the job coordination, the `DocumentGenerationManager` assigns jobs to the `Generators` in groups of more than one job that are part of the same batch. If a `Generator` fails to complete its jobs, the `DocumentGenerationManager` can restart these failed jobs.
It prioritizes jobs based on their deadlines and schedules them according to *P1*.
- **Super-component:** None
- **Sub-components:** Completer, GenerationManager, KeyCache, Scheduler, TemplateCache

Provided interfaces

- InsertJobs
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
- NotifyCompleted
 - `void notifyCompletedAndGiveMeMore(GeneratorId id)`
 - * Effect: The `DocumentGenerationManager` gets notified that the document processing jobs assigned to the `Generator` identified by an `id` are completed.
 - * Exceptions: None
 - `void notifyCompletedAndIAMShuttingDown(GeneratorId id)`
 - * Effect: The `DocumentGenerationManager` gets notified that the document processing jobs assigned to the `Generator` identified by an `id` are completed.
 - * Exceptions: None

A.11 DocumentStorageCache

- **Description:** The `DocumentStorageCache` is responsible for storing the `DocumentIds` and `UserIds` when the PDSDB fails. According to Av2, the system should temporarily store at least 3 hours of documents to be delivered via the personal document store. When the PDSDB fails, the documents that are supposed to also be saved in the PDSDB are saved in the `DocumentDB`, just as usual. But in this case the `DocumentStorageManager` also stores the `DocumentIds` and `UserIds` of those documents in the `DocumentStorageCache` for at least 3 hours. This way, the `DocumentStorageManager` can transfer these documents from the `DocumentDB` to the PDSDB using this information if the PDSDB comes back online within 3 hours. The requirements do not specify what happens after 3 hours, so in this architecture, the behaviour after those 3 hours is undefined.
- **Super-component:** `DocumentStorageFunctionality`
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - `SomeException`: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.12 DocumentStorageFunctionality

- **Description:** The `DocumentStorageManager` is responsible for storing the generated documents in the correct database. If the document belongs to a registered recipient, it sends a write request to both the `DocumentDB` and the PDSDB. Otherwise, it only sends a write request to the `DocumentDB`, which stores all the documents.
It is also responsible for copying documents from the `DocumentDB` to the PDSDB when an Unregistered Recipient registers to the eDocs system.
Another responsibility of the `DeliveryFunctionality` is storing at least 3 hours of documents when the PDSDB fails.
- **Super-component:** None
- **Sub-components:** `DocumentStorageCache` and `DocumentStorageManager`

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - `SomeException`: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation

- * Exceptions: None
- InterfaceB
 - returnType2 operation3()
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.13 DocumentStorageManager

- **Description:** The `DocumentStorageManager` is responsible for storing the generated documents in the correct database. If the document belongs to a registered recipient, it sends a write request to both the `DocumentDB` and the `PDSDB`. Otherwise, it only sends a write request to the `DocumentDB`, which stores all the documents.
It is also responsible for copying documents from the `DocumentDB` to the `PDSDB` when an Unregistered Recipient registers to the eDocs system.
Another responsibility of the `DocumentStorageManager` is storing at least 3 hours of documents when the `PDSDB` fails.
- **Super-component:** `DocumentStoragefunctionlity`
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - returnType1 operation1(ParamType param) throws SomeException
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
- InterfaceB
 - returnType2 operation3()
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.14 DeliveryFunctionality

- **Description:** The `DeliveryFunctionality` is responsible for delivering the documents generated by eDocs to the recipients.
- **Super-component:** `ChannelDispatcher`, `EmailFacade`, `Print&PostalServiceFacade`, `ZoomitFacade` and `ZoomitDeliveryCache`
- **Sub-components:** the direct sub-components, if any.

Provided interfaces

- InterfaceA
 - returnType1 operation1(ParamType param) throws SomeException
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
- InterfaceB
 - returnType2 operation3()
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.15 EDocsAdminClient

- **Description:** The `EDocsAdminClient` is external to the eDocs system and represents a client device of an administrator of eDocs that communicates with the eDocs System.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - `SomeException`: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.16 EmailChannel

- **Description:** The `EmailChannel` is responsible for the delivery emails. It is external to the eDocs system and represents a mail server of an e-mail provider.
- **Super-component:** None
- **Sub-components:** None.

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - `SomeException`: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.17 EmailFacade

- **Description:** The **EmailFacade** is responsible for creating and sending emails used in the delivery of documents. It will send documents to Unregistered recipients by e-mail when receipt tracking is turned off. When receipt tracking is turned on for an Unregistered Recipient, it will send an e-mail containing a short description of the received document and a unique link, which can be followed to get document. For Registered Recipients, it will send an e-mail containing a short description of the document and a link to the document. It also marks jobs as sent using the **JobManager**.
- **Super-component:** **DeliveryFunctionality**
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - **SomeException:** Describe when the exception is thrown.
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.18 Generator

- **Description:** A **Generator** generates the documents and forwards them to **DeliveryFunctionality** to store and deliver them. Its availability is monitored by the **DocumentGenerationManager** with the **Ping** interface. A **Generator** is also responsible of notifying the **NotificationHandler** that it does not have all of the data required to fill in the template.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- AssignJobs
 - `void assignJobs(CompletePartialBatchData batchData)`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - **SomeException:** Describe when the exception is thrown.
- Startup/ShutDown
 - `void startUp(GeneratorId generatorId)`
 - * Effect: Starts up the **Generator** instance and gives it the given **GeneratorId**.
 - * Exceptions: None
 - `void shutDown()`
 - * Effect: The **Generator** completes its assigned group of document generation jobs and report back completion to the **DocumentGenerationManager**, after which it shuts down.
 - * Exceptions: None
- Ping
 - `Echo ping()`
 - * Effect: The **Generator** will respond to the ping request by sending an echo response. This is used by the **GeneratorManager** to check whether the **Generator** is available.
 - * Exceptions: None

A.19 GeneratorManager

- **Description:** The `GeneratorManager` is responsible for monitoring the `Generator` instances. It starts up or shuts down these instances based on the number of required instances indicated by the `Scheduler`.
- **Super-component:** `DocumentGenerationManager`
- **Sub-components:** None

Provided interfaces

- `NotifyCompleted`
 - `void notifyCompletedAndGiveMeMore(GeneratorId id)`
 - * Effect: The `DocumentGenerationManager` gets notified that the document processing jobs assigned to the `Generator` identified by an `id` are completed.
 - * Exceptions: None

A.20 JobDBShard

- **Description:** The `JobDBShard` is responsible for storing partition of all the jobs.
- **Super-component:** `JobManager`
- **Sub-components:** None

Provided interfaces

- `InterfaceA`
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - `SomeException`: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- `InterfaceB`
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.21 JobDBShardingManager

- **Description:** The `JobDBShardingManager` manages the storage of the documents over multiple `JobDBShards`.
- **Super-component:** `JobManager`
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.22 JobFacade

- **Description:** The JobFacade is responsible creating jobs and storing them using the JobDBShadingManager over different JobDBShards. It is responsible for retrieving data connected to a specific job. It can retrieve the raw data or customer organization info for specific jobs. The JobFacade is also used for marking jobs as sent and received.
- **Super-component:** JobManager
- **Sub-components:** None

Provided interfaces

- SetStatus
 - `void setJobStatusAsTemporarilyFailed(List<JobId> statusesOfJobs)`
 - * Effect: The JobManager marks the job as “temporarily failed” for each of the jobs identified by the given JobIds. Used by the DocumentGenerationManager for jobs that were assigned to a failed Generator instance.
 - * Exceptions: None
- GetBatchData
 - `Tuple<JobId, RawData> getRawData(List<JobId> jobIds)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
 - `BatchMetaData getMetaData(BatchId batchId)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- JobMgmt
 - `void createJobsAndInitiateProcessing(Customerid cuid, TimeStamp deadline, List<rawdataId> rawDataIds, batchId batchId)`
 - * Effect: The JobFacade creates a job for each of the raw data identifiers it gets as an argument. The job contains the time of deadline for the document processing job and the batch identifier. It stores the jobs and inserts the jobs in the Scheduler, together with their deadline and batch identifier.
 - * Exceptions: None

A.23 JobManager

- **Description:** The `JobManager` is responsible creating jobs and storing them. It is responsible for retrieving data connected to a specific job. It can retrieve the raw data or customer organization info for specific jobs. The `JobFacade` is also used for marking jobs as sent and received.
- **Super-component:** None
- **Sub-components:** `JobFacade`, `JobDBShardingManager` and `JobDBShard`

Provided interfaces

- `SetStatus`
 - `void setJobStatusAsTemporarilyFailed(List<JobId> statusesOfJobs)`
 - * Effect: The `JobManager` marks the job as “temporarily failed” for each of the jobs identified by the given `JobIds`. Used by the `DocumentGenerationManager` for jobs that were assigned to a failed `Generator` instance.
 - * Exceptions: None
- `GetBatchData`
 - `Tuple<JobId, RawData> getRawData(List<JobId> jobIds)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
 - `BatchMetadata getMetadata(BatchId batchId)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.24 KeyCache

- **Description:** The `KeyCache` caches the keys which are most recently used for document generation. The `Completer` has to fetch a key every time a `Generator` instance requests new jobs, while the key will be the same for all jobs belonging to the same batch. The `KeyCache` avoids that the key storage system becomes a bottleneck for document generations. The keys are cached based on the `CustomerId` of a Customer Organization.
- **Super-component:** `DocumentGenerationManager`.
- **Sub-components:** None

Provided interfaces

- `GetKey`
 - `Key getKey(CustomerId customerId)` throws `NoSuchKeyException`
 - * Effect: The `KeyCache` looks into its cache for the `Key` belonging to the customer organisation with id `customerId`. If the `Key` is in its cache, it returns it. If the `Key` is not in its cache, it asks `OtherDB` for the `Key` and stores it in its cache, after which it returns that `Key`.
 - * Exceptions:
 - `NoSuchKeyException`: Thrown if there is no key for the given `customerId`.

A.25 LinkMappingDB

- **Description:** The `LinkMappingDB` is responsible for actually storing a mapping between unique links and `DocumentIds`. With this information, it also stores information about where a document can be found, i.e. only in the `DocumentDB` or both in the `PDSDB` and the `DocumentDB`.
- **Super-component:** `LinkMappingFunctionality`
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.26 LinkMappingManager

- **Description:** The `LinkMappingManager` is responsible for creating unique links which points to a document. It also sends read and write requests to the `LinkMappingDB` to get and store the mappings between the unique links and the documents.
- **Super-component:** `LinkMappingFunctionality`
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.27 LinkMappingFunctionality

- **Description:** The `LinkMappingFunctionality` is responsible for creating unique links which point to documents. It is also responsible for storing these mappings and ultimately mapping a link to a document.
- **Super-component:** None
- **Sub-components:** `LinkMappingManager` and `LinkMappingFunctionality`

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.28 PDSDB

- **Description:** The PDSDB component is responsible for storing the database of documents in the personal document stores. That database is separated from all other persistent data so that its failure *“does not affect the availability of other types of persistent data”*, as required by *Av2*.
- **Super-component:** None
- **Sub-components:** PDSDBReplica, PDSLongTermDocumentManager, PDSReplicationManager

Provided interfaces

- DocumentMgmt
 - `Tuple<Document, Metadata> getDocument(DocumentId id)`
 - * Effect: The PDSDB will fetch and return the document corresponding to `DocumentId` `id`.
 - * Exceptions: None
 - `List<Document> getAllDocumentMetaDataOf(RecipientId recipientId)`
 - * Effect: The PDSDB will fetch and return all the meta-data of the documents belonging to the Registered Recipient identified by `recipientId`.
 - * Exceptions: None
 - `void storeDocument(DocumentId id, Document doc, DocumentMetaData md)`
 - * Effect: The PDSDB will store the given document `doc` together with the provided meta-data `md`.
 - * Exceptions: None
 - `void storeDocuments(List<Tuple<DocumentId, Document, DocumentMetaData>> documentList)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
 - `List<Tuple<DocumentId, DocumentMetaData>> getAllDocumentMetaData(RecipientId recipientId) throws PDSUnavailableException`
 - * Effect: The PDSDB fetches and returns the meta-data of all the documents of the Registered Recipient identified by `recipientId`.
 - * Exceptions:
 - PDSUnavailableException: Thrown if the personal document store is unavailable.

A.29 PDSDBReplica

- **Description:** The PDSDBReplica is responsible for actually storing the documents.
- **Super-component:** PDSDB
- **Sub-components:** None

Provided interfaces

- ExtendedDocumentMgmt
 - `List<Document> getAllDocumentsOf(TimeStamp whenFailed)`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `List<Document> getDocumentsSince(TimeStamp whenFailed)`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void storeDocuments(List<Tuple<DocumentId, Document, DocumentMetaData>> documentList)`
 - * Effect: The PDSDBReplica will store the documents and their meta-data.
 - * Exceptions: None
 - `void storeDocument(DocumentId documentId, Document document, DocumentMetaData md)`
 - * Effect: The PDSDBReplica stores the given document with its DocumentId and meta-data.
 - * Exceptions: None
 - `Tuple<Document, MetaData> getDocument(DocumentId id)`
 - * Effect: The PDSDB will fetch and return the document corresponding to DocumentId id.
 - * Exceptions: None
 - `List<Tuple<DocumentId, DocumentMetaData>> getAllDocumentMetaData(RecipientId recipientId)`
`throws PDSUnavailableException`
 - * Effect: The PDSDBReplica fetches and returns the meta-data of all the documents of the Registered Recipient identified by recipientId.
 - * Exceptions: None
- Ping
 - `Echo ping()`
 - * Effect: The PDSDBReplica will respond to the ping request by sending an echo response. This is used by the PDSReplicationManager to check whether the PDSDBReplica is available.
 - * Exceptions: None

A.30 PDSFacade

- **Description:** The PDSFacade is responsible for handling all read requests from Registered recipients to the PDSDB. It checks whether the documents requested by a Registered Recipient actually belong to him or her.
It is also responsible for handling queries when Registered Recipients search for documents in their personal document store. It fetches the meta-data of the documents of the searching Registered Recipients, filters through that meta-data and returns an overview of the document meta-data and document identifiers matching the search criteria.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- PDSDBDocMgmt

- `List<Tuple<DocumentId, DocumentMetaData>> getAllDocumentMetaData(RecipientId recipientId)`
throws `PDSUnavailableException`
 - * Effect: The `PDSFacade` fetches and returns the meta-data of all the documents of the Registered Recipient identified by `recipientId`.
 - * Exceptions:
 - `PDSUnavailableException`: Thrown if the personal document store is unavailable.
- `Tuple<Document, DocumentMetaData> getDocument(DocumentId documentId, RecipientId recipientId)`
throws `PDSUnavailableException`, `DocumentUnavailableException`
 - * Effect: The `DocumentFacade` first checks whether the document identified by the given document identifier actually belongs to the registered recipient. If it does, it fetches the document from the PDSDB and returns it. Otherwise, it throws an exception.
 - * Exceptions:
 - `PDSUnavailableException`: Thrown if the personal document store is unavailable.
 - `DocumentUnavailableException`: Thrown if the given document is not available for the given user. An example is when the document identified by the `documentId` exists, but the document does not belong to the Registered Recipient. Another example is when the document identifier does not exist.

A.31 PDSLongTermDocumentManager

- **Description:** The `PDSLongTermDocumentManager` is responsible for managing the different storage clusters. Each cluster consists of a `PDSReplicationManager` and one or more `PDSDBReplica` instances. In the architecture, two clusters are defined. The `PDSLongTermDocumentManager` reads to and write from clusters, and periodically transfers documents from the one cluster to the other.
- **Super-component:** PDSDB
- **Sub-components:** None

Provided interfaces

- DocumentMgmt

- `Tuple<Document, MetaData> getDocument(DocumentId id)`
 - * Effect: The PDSDB will fetch and return the document corresponding to `DocumentId id`.
 - * Exceptions: None
- `List<Document> getAllDocumentMetaDataOf(RecipientId recipientId)`
 - * Effect: The `PDSLongTermDocumentManager` will fetch and return all the meta-data of the documents belonging to the Registered Recipient identified by `recipientId`.
 - * Exceptions: None
- `void storeDocument(DocumentId id, Document doc, DocumentMetaData md)`
 - * Effect: The PDSDB will store the given `documentdoc` together with the provided meta-data `md`.
 - * Exceptions: None
- `void storeDocuments(List<Tuple<DocumentId, Document, DocumentMetaData>> documentList)`
 - * Effect: The `PDSLongTermDocumentManager` will store the documents and their meta-data.
 - * Exceptions: None
- `List<Tuple<DocumentId, DocumentMetaData>> getAllDocumentMetaData(RecipientId recipientId)`
throws `PDSUnavailableException`
 - * Effect: The `PDSLongTermDocumentManager` fetches and returns the meta-data of all the documents of the Registered Recipient identified by `recipientId`.
 - * Exceptions:
 - `PDSUnavailableException`: Thrown if the personal document store is unavailable.

A.32 PDSReplicationManager

- **Description:** The PDSReplicationManager is responsible for managing the PDSDBReplicas. The PDSReplicationManager passes read requests to one PDSDBReplica and writes to all PDSDBReplicas. It monitors their availability using the ping/echo.
- **Super-component:** PDSDB
- **Sub-components:** None

Provided interfaces

- ExtendedDocumentMgmt
 - Tuple<Document, Metadata> getDocument(DocumentId documentId)
 - * Effect: The PDSReplicationManager will fetch and return the document corresponding to DocumentId id.
 - * Exceptions: None
 - List<Tuple<DocumentId, DocumentMetadata>> getAllDocumentMetadata(RecipientId recipientId) throws PDSUnavailableException
 - * Effect: The PDSReplicationManager fetches and returns the meta-data of all the documents of the Registered Recipient identified by recipientId.
 - * Exceptions:
 - PDSUnavailableException: Thrown if the personal document store is unavailable.
 - void storeDocument(DocumentId id, Document doc, DocumentMetadata md)
 - * Effect: The PDSReplicationManager will store the given document doc together with the provided meta-data md.
 - * Exceptions: None
 - void storeDocuments(List<Tuple<DocumentId, Document, DocumentMetadata>> documentList)
 - * Effect: The PDSReplicationManager will store the given list of documents and their meta-data.
 - * Exceptions: None

A.33 Print&PostalServiceChannel

- **Description:** The Print&PostalServiceChannel is responsible for the printing the document and sending it by mail. It is external to the eDocs system and represents the servers of a print & postal service.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - returnType1 operation1(ParamType param) throws SomeException
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - void operation2(ParamType2 param)
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - returnType2 operation3()
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.34 Print&PostalServiceFacade

- **Description:** The `Print&PostalServiceFacade` is responsible for delivering a document to the `Print&PostalChannel` so it can be printed and sent by mail. It also marks jobs as sent using the `JobManager`.
- **Super-component:** `DeliveryFunctionality`
- **Sub-components:** None

Provided interfaces

- `InterfaceA`
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - `SomeException`: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- `InterfaceB`
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.35 NotificationHandler

- **Description:** The `NotificationHandler` is responsible for sending notifications to the appropriate parties, e.g. the eDocs operators and the customer administrators.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- `NotifyOperator`
 - `void notifyOperatorOfPDSDBReplicaFailure(PDSDBReplicaId replicaId,TimeStamp dateTime)`
 - * Effect: The `NotificationHandler` will send the given `PDSDBReplicaId` of the failed `PDSDBReplica` with the given time of failure `dateTime` to the eDocs operators. This method is called by a `PDSReplicationManager`.
 - * Exceptions: None
 - `void notifyOperatorOfDocumentGenerationFailure(NotificationMessage msg,TimeStamp whenFailed)`
 - * Effect: The `NotificationHandler` will send a textual message `msg` to the eDocs operators, which contains further information about the specific failure. This method is called by the `DocumentGenerationManager`. More specifically, it is called by the `GeneratorManager`.
 - * Exceptions: None

A.36 OtherDB

- **Description:** The `OtherDB` is responsible for storing all information that is not required to be stored separately by non-functional requirements. For example, it stores the raw data and data about customer organizations and registered recipients. It also stores the templates for documents and the keys of customer organizations to sign the documents during generation.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- **GetKey**
 - Key `getKey(CustomerId customerId)`
 - * Effect: The `OtherDB` returns the key belonging to the Customer Organization identified by `customerId`.
 - * Exceptions:
 - `NoSuchKeyException`: Thrown if there is no key for the given `customerId`.
- **TemplateMgmt**
 - Template `getTemplate(CustomerId customerId, DocumentType documentType, TimeStamp whenReceived)`
 - * Effect: The `OtherDB` returns the Template belonging to the customer organisation with id `customerId` corresponding to a document of type `documentType` and received at time `whenReceived`.
 - * Exceptions:
 - `NoSuchTemplateException`: Thrown if there is no template for the given arguments.
 - Boolean `storeDocumentTemplate(CustomerId cuId, DocumentType documentType, Template template, TimeStamp whenReceived)` throws `InvalidDocumentTypeException`
 - * Effect: The `OtherDB` stores the given template with the given time stamp for the customer organization identified by the given `CustomerId`.
 - * Exceptions:
 - `InvalidDocumentTypeException`: Thrown if the given document type is invalid or not allowed for the given customer organization.
- **UserDataMgmt**
 - Credentials `getRegisteredRecipientCredentials(RecipientId recipientId)` throws `NoSuchRecipientException`
 - * Effect: The `OtherDB` returns the credentials belonging to the Registered Recipient identified by `recipientId`.
 - * Exceptions:
 - `NoSuchRecipientException`: Thrown if no Registered Recipient with the given credentials exists.
 - Credentials `getCustomerOrganizationCredentials(CustomerId customerId)` throws `NoSuchCustomerOrganizationException`
 - * Effect: The `OtherDB` returns the credentials belonging to the Customer Organization identified by `customerId`.
 - * Exceptions:
 - `NoSuchCustomerOrganizationException`: Thrown if no Customer Organization with the given credentials exists.
 - List<Tuple<DocumentType, TimeStamp>> `getPossibleDocumentTypes(CustomerId customerId)`
 - * Effect: The `OtherDB` returns the a list of document types that the Customer Organization identified by `customerId` can generate. For each document type, it also returns the time when the last Template for that document type was uploaded.
 - * Exceptions: None
 - Boolean `isAllowedDocumentType(CustomerId cuid, DocumentType docType)`
 - * Effect: The `OtherDB` checks in the SLA of the customer organization with customer identifier `cuid` if the document type `docType` is allowed. It returns true if it is allowed, otherwise it returns false.
 - * Exceptions: None
 - void `indicateWhichRecurringBatch(CustomerId cuid, BatchMetaData batchMetaData)`

- * Effect: The `OtherDB` stores to which of the recurring batches of the customer organisation identified by `cuid` this batch corresponds. The `batchMetaData` contains a description of which recurring batch this batch is.
 - * Exceptions: None
- `int getNumberOfDocumentsMaximallyAllowed(CustomerId cuid, BatchMetaData batchMetaData)`
 - * Effect: The `OtherDB` returns the number of documents that is maximally allowed for the customer organization identified by `cuid` for this recurring batch. The `batchMetaData` contains a description of which recurring batch this batch is.
 - * Exceptions: None
- `BatchDataMgmt`
 - `BatchId storeBatchMetaData(BatchMetaData batchMetaData)`
 - * Effect: The `OtherDB` stores the given `batchMetaData` and returns the batch identifier with which this meta-data can be queried.
 - * Exceptions: None
 - `BatchMetaData getBatchMetaData(BatchId batchId)`
 - * Effect: The `OtherDB` returns the meta-data if the batch identified by the given `batchId`.
 - * Exceptions: None
- `RawDataMgmt`
 - `List<RawDataId> storeRawData(List<RawData> rawDataEntries)`
 - * Effect: The `OtherDB` stores the given raw data entries. It returns the identifiers of the stored raw data entries.
 - * Exceptions: None
 - `List<RawData> getRawData(List<RawDataId> rawDataIds)`
 - * Effect: The `OtherDB` returns the raw data entries corresponding to the given raw data identifiers.
 - * Exceptions: None

A.37 RecipientClient

- **Description:** The `RecipientClient` is external to the eDocs system and represents a client device of an unregistered or registered recipient of eDocs that communicates with the eDocs System.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- `InterfaceA`
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - `SomeException`: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- `InterfaceB`
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.38 RawDataHandler

- **Description:** The `RawDataHandler` is responsible for verifying the raw data and its entries. It forwards the validated raw data to the `JobManager` to create jobs.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- `InputRawData`
 - `List<RawData> initiateDocumentProcessing(CustomerId coid, DocumentType docType, RawDataPacket rawDataPacket, DocumentPriority priority, Boolean isRecurring, TimeStamp isReceived)` throws `InvalidDocumentTypeException`, `InvalidRawDataPackageException`, `ToManyRawDataEntriesException`
 - * Effect: The `RawDataHandler` verifies the received raw data. If it is correct, it stores the raw data and batch meta-data in the `OtherDB`. Next, it calculates the deadline of the document processing. It generates a `BatchMetaData` object containing information about the batch and stores it in the `OtherDB`. It also stores the valid raw data entries in the `OtherDB`. The `RawDataHandler` then informs the `JobFacade` to create jobs and to initiate document processing for the valid raw data entries. It returns a list of the invalid raw data entries.
 - * Exceptions:
 - `InvalidDocumentTypeException`: Thrown if the given document type is invalid or not allowed for the given customer organization.
 - `InvalidRawDataPackageException`: Thrown if the raw data package contains an error (e.g. an invalid Excel file or incorrectly formatted XML).
 - `ToManyRawDataEntriesException`: Thrown if the batch is a recurring batch and contains more raw data entries than maximally allowed for this batch.

A.39 RecipientFacade

- **Description:** The `RecipientFacade` is responsible for the interaction of Registered and Unregistered Recipients with the eDocs system. It provides methods for authentication, for consulting the personal document store, for downloading documents, ...
- **Super-component:** `Userfunctionality`
- **Sub-components:** None

Provided interfaces

- `AuthN`
 - `SessionId login(Credentials credentials)` throws `InvalidCredentialsException`
 - * Effect: The `RecipientFacade` forwards the given `credentials` to the `AuthenticationHandler`, which verifies them and returns a new session identifier if correct. This session identifier can be used in future requests to the `RecipientFacade`.
 - * Exceptions:
 - `InvalidCredentialsException`: Thrown if the `AuthenticationHandler` indicated that the given credentials were incorrect.
 - `Boolean logout(SessionId sessionId)`
 - * Effect: The `RecipientFacade` removes the session corresponding to the `sessionId` using the `AuthenticationHandler`. As a result, this session cannot be used anymore to access the system without logging in again. If no session corresponds to the `sessionId`, it does not exist, nothing is changed but no exception is thrown.
 - * Exceptions: None

- **DocumentMgmt**
 - `PDSOverview getPDSOverview(SessionId, RecipientId recipientId)`
throws `NotAuthenticatedException`, `PDSUnavailableException`
 - * **Effect:** The `RecipientFacade` first verifies the given session identifier `sessionId` using the `AuthenticationHandler`. The `RecipientFacade` then requests all the document meta-data of the documents of the recipient identified by `recipientId` from the `PDSFacade`. It generates a document overview , which is returned to the caller.
 - * **Exceptions:**
 - `NotAuthenticatedException`: Thrown if the given session identifier is invalid.
 - `PDSUnavailableException`: Thrown if the personal document store is unavailable.
 - `Tuple<Document, DocumentMetaData> lookupDocument(RecipientId recipientId, SessionId sessionId, DocumentId documentId)` throws `NotAuthenticatedException`
 - * **Effect:** The `RecipientFacade` first verifies the given session identifier `sessionId` using the `AuthenticationHandler`. The `RecipientFacade` then gets the document identified by `documentId` and its meta-data using the `PDSFacade`, which checks whether the document belongs to the Registered Recipient, and returns them.
 - * **Exceptions:**
 - `NotAuthenticatedException`: Thrown if the given session identifier is invalid.
 - `DocumentUnavailableException`: Thrown if the given document is not available for the given user. An example is when the document identified by the `documentId` exists, but the document does not belong to the Registered Recipient. Another example is when the document identifier does not exist.

A.40 RegistrationManager

- **Description:** The `RegistrationManager` is responsible for the registration or unregistered recipients and customer organizations. For the registration of unregistered recipients, the `RegistrationManager` gets called by the `RecipientFacade`, as recipients can register themselves. Customer organizations get registered by an eDocs operator, so the `EDocsAdminClient` calls those methods.
- **Super-component:** `UserFunctionality`
- **Sub-components:** None

Provided interfaces

- **InterfaceA**
 - `returnType1 operation1(ParamType param)` throws `SomeException`
 - * **Effect:** Describe the effect of the operation
 - * **Exceptions:**
 - `SomeException`: Describe when the exception is thrown.
- **InterfaceB**
 - `returnType2 operation3()`
 - * **Effect:** Describe the effect of the operation
 - * **Exceptions:** None

A.41 Scheduler

- **Description:** The `Scheduler` receives the new jobs initiated by a Customer Organization and adds them to a queue of all jobs that have not been processed yet. To lower the size of this queue, the `Scheduler` is only given the information it needs, i.e., the id of the batch, its deadline and the ids of the individual jobs. The raw data of each job and the meta-data of the batch is stored in `OtherDB` and fetched by the `Completer` when needed. The `Scheduler` also indicates to the `GenerationManager` the number of required `Generator` instances through its `GetStatistics` interface.

- **Super-component:** DocumentGenerationManager
- **Sub-components:** None

Provided interfaces

- GetNextJobs
 - `Tuple<BatchId, List<JobId>> getNextJobs()`
 - * Effect: The **Scheduler** returns the **JobIds** of the group of jobs that belong to the batch identified by **BatchId** that should be generated next. This method is called by the **GeneratorManager** when a **Generator** instance requires a new group of jobs.
 - * Exceptions: None
 - `Tuple<BatchId, List<JobId>> jobsCompletedAndGiveMeMore(List<JobId>)`
 - Effect: The **Scheduler** gets notified that the document processing jobs belonging to the list of **JobIds** are completed. It returns the a list of **JobIds** belonging to a batch identified by **BatchId**. The returned list of **JobIds** identify document processing jobs which are not yet started.
 - Exceptions: None
- InsertJobs
 - `void insertJobs(BatchId batchId,TimeStamp deadline, List<JobId> jobIds)`
 - * Effect: The **Scheduler** adds the jobs identified by their **JobId** to its queue of all jobs that have not been processed yet. To lower the size of this queue, the Scheduler is only given the information it needs, i.e., the id of the batch, its deadline and the ids of the individual jobs. This method provides new jobs synchronously to the **Scheduler**, which it schedules synchronously. This means that when the method call returns, the given jobs are scheduled.
 - * Exceptions: None
- GetStatistics
 - `int getNumberOfFutureJobs()`
 - * Effect: The **Scheduler** returns the amount of documents that should be generated in the near future. The **GeneratorManager** queries this method at regular intervals and adjusts the number of **Generator** instances accordingly.
 - * Exceptions: None

A.42 SessionDB

- **Description:** The **SessionDB** stores the session identifiers for currently active sessions.
- **Super-component:** UserFunctionality
- **Sub-components:** None.

Provided interfaces

- SessionMgmt
 - `RecipientId getRecipientId(SessionId sessionId) throws NoSuchSessionException`
 - * Effect: The **SessionDB** fetches and returns the Registered Recipient's identifier corresponding to the **sessionId** from the **sessionDB**.
 - * Exceptions:
 - **NoSuchSessionException**: Thrown if no session exists with the given identifiers, or if the session belongs to a customer organization.
 - `CustomerId getCustomerId(SessionId sessionId) throws NoSuchSessionException`
 - * Effect: The **SessionDB** fetches and returns the Customer Organization's identifier corresponding to the **sessionId** from the **sessionDB**.

- * Exceptions:
 - `NoSuchSessionException`: Thrown if no session exists with the given identifier, or if the session belongs to a registered recipient.
- `SessionId openSession(RecipientId recipientId)`
 - * Effect: The `SessionDB` generates a new session identifier for the given `recipientId` and stores this as an active session.
 - * Exceptions: None
- `void closeSession(SessionId sessionId) throws NoSuchSessionException`
 - * Effect: The `SessionDB` closes the active session associated with the given `sessionId`.
 - * Exceptions:
 - `NoSuchSessionException`: Thrown if no session exists with the given identifier.
- `Map<SessionAttributeKey, SessionAttributeValue> isValidSession(SessionId sessionId) throws NoSuchSessionException`
 - * Effect: The `SessionDB` verifies whether a session with the given id exists in the `SessionDB` and if so, returns all its associated attributes.
 - * Exceptions:
 - `NoSuchSessionException`: Thrown if no session exists with the given identifiers.

A.43 TemplateCache

- **Description:** The `TemplateCache` caches the templates which are most recently used for document generation. The `Completer` has to fetch a template every time a `Generator` instance requests new jobs, while the template will be the same for all jobs belonging to the same batch. The `TemplateCache` avoids that the template storage system becomes a bottleneck for document generations. The templates are cached based on the `CustomerId` of a Customer Organization, the type of the document and the date and time at which the batch was provided by the Customer Organization (in order to account for template updates).
- **Super-component:** `DocumentGenerationManager`
- **Sub-components:** None

Provided interfaces

- `GetTemplate`
 - `Template getTemplate(CustomerId customerId, DocumentType documentType, Timestamp whenReceived)`
 - * Effect: The `TemplateCache` looks into its cache for the `Template` belonging to the customer organisation with id `customerId` corresponding to a document of type `documentType` and received at time `whenReceived`. If the `Template` is in its cache, it returns it. If the `Template` is not in its cache, it asks `OtherDB` for the `Template` and stores it in its cache, after which it returns that `Template`.
 - * Exceptions:
 - `NoSuchTemplateException`: Thrown if there is no template for the given arguments.

A.44 UserFunctionality

- **Description:** The `UserFunctionality` is responsible for the interaction of registered recipients, unregistered recipients, customer organizations and eDocs operators with the eDocs system. It provides methods to register, to login and to logout, to consult the personal document store and download documents, to consult the status of document processing jobs, ...
- **Super-component:** None
- **Sub-components:** `RecipientFacade`, `CustomerOrganizationClient`, `EDocsadminfacade`, `RegistrationManager`, `AuthenticationHandler` and `SessionDB`

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.45 ZoomitChannel

- **Description:** The `ZoomitChannel` is responsible for delivering documents via Zoomit. It is external to the system and represents the servers of Zoomit to which a document can be sent.
- **Super-component:** None
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.
 - `void operation2(ParamType2 param)`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None
- InterfaceB
 - `returnType2 operation3()`
 - * Effect: Describe the effect of the operation
 - * Exceptions: None

A.46 ZoomitFacade

- **Description:** The `ZoomitFacade` is responsible for sending documents to Unregistered Recipients through Zoomit. It is also responsible for receiving messages from Zoomit when a document has been received by Zoomit or when a Zoomit user has received his or her document. The `ZoomitFacade` can use the `JobManager` to mark jobs as sent or received.
- **Super-component:** `DeliveryFunctionality`
- **Sub-components:** None

Provided interfaces

- InterfaceA
 - `returnType1 operation1(ParamType param) throws SomeException`
 - * Effect: Describe the effect of the operation
 - * Exceptions:
 - SomeException: Describe when the exception is thrown.

B Defined data types

List and describe all data types defined in your interface specifications. List them alphabetically for ease of navigation.

- **BatchId**: A piece of data uniquely identifying a batch of document processing jobs in the system. This architecture does not specify the exact format of this identifier, but possibilities are a long integer, a string, a URL etc.
- **BatchMetaData**: A data structure listing the metadata belonging to a batch of jobs. This includes the **CustomerId** of a Customer Organization, the **DocumentType** of the documents to be generated, the **TimeStamp** of when the batch was received, ...
- **CompletePartialBatchData**: A complex data structure listing all data a **Generator** needs to complete document generation jobs that are part of the same batch. It contains an array of **Tuple<JobId, RawData>**. The **JobIds** identify jobs that are all part of the same batch. The **RawData** belongs to these document processing jobs. Also listed in the **BatchMetaData** are the values of the **BatchMetaData**, **Key** and **Template** data types belonging to the batch.

CompletePartialBatchData also contains a **BatchMetaData** entry, a **Key** and a **Template**. *Important to note*: a value of **CompletePartialBatchData** contains all information necessary to generate **some** jobs of belonging to same batch. It does not have to contain the information of all jobs belonging to same batch.
- **Credentials**: The authentication credentials of a Registered Recipient or Customer Organization. The credentials always contain an identifier of the recipient or customer organization and a proof of his or her identity. The architecture does not specify the specific credentials used, but a possibility is using a username and password.
- **CustomerId**: A piece of data uniquely identifying a Customer Organization in the system. This architecture does not specify the exact format of this identifier, but possibilities are a long integer, a string, a URL etc.
- **DeliveryInfo**:
- **Document**: A data file corresponding to a document. The architecture specifies the format of this data type as a PDF-file.
- **DocumentId**: A piece of data uniquely identifying a document in the system.
- **DocumentMetaData**:
- **DocumentPriority**: A data type representing the priority of document generation jobs. They have values representing the Critical, Diamond, Gold and Silver priorities. The exact format of this data type is not specified by the architecture.
- **DocumentType**: A piece of data describing the type of a document. This architecture does not specify the exact format of this data type, but possibilities are a long integer, a string, a URL etc.
- **Echo**: The response to a ping message. This data element does not contain any meaningful data.
- **Error**: Description of data type.

- **GeneratorId**: A piece of data uniquely identifying a **Generator** in the system. This architecture does not specify the exact format of this identifier, but possibilities are a long integer, a string, a URL etc.
- **JobBatch**: Description of data type.
- **JobId**: A piece of data uniquely identifying a document processing job in the system.
- **Key**: A data structure containing the key of the Customer Organization which is used to sign its documents during the generation process. This architecture does not specify the exact format of this data type, but possibilities are a long integer, a string, a URL etc.
- **NotificationMessage**: A textual message which can be used to include extra information about the event of the notification.
- **PDSOverview**: The overview of a personal document store that can be shown to a Registered Recipient. Through this overview, the Registered Recipient can consult documents. The architecture does not specify the exact format of such an overview, but a likely possibility is an HTML page.
- **PDSDBReplicaId**: A piece of data uniquely identifying a **PDSDBReplica** in the system. This architecture does not specify the exact format of this identifier, but possibilities are a long integer, a string, a URL etc.
- **RawData**: A data structure listing a raw data entry used in a document processing job.
- **RawDataPacket**: A data structure containing all the raw data entries in a batch, given by the customer organization to the **CustomerOrganizationFacade**. The architecture does not specify this data structure, but possibilities are an XML file or Excel file.
- **RawDataId**: A piece of data uniquely identifying a raw data entry in the system. This architecture does not specify the exact format of this identifier, but possibilities are a long integer, a string, a URL etc.
- **RecipientId**: A piece of data uniquely identifying a Registered Recipient in the system. This architecture does not specify the exact format of this identifier, but possibilities are a long integer, a string, a URL etc.
- **SessionId**: A piece of data uniquely identifying a session of a registered recipient of customer organization in the eDocs system. This contains at least the user identifier (i.e. the **CustomerId** for a customer organization or the **RecipientId** for a registered recipient) and the time the session was initiated.
- **SessionAttributeKey**: The key of an attribute attached to a session. This architecture does not specify the exact format of this key. a possible value is a long integer or a flat string.
- **SessionAttributeValue**: The value of an attribute attached to a session. This value can be of any primitive type.
- **TimeStamp**: The representation of a time (i.e. date and time of day) in the system.
- **Template**: A document used as a template for the generation of documents.
- **TemplateId**: A data structure uniquely identifying a template in the system. It lists three values. It contains **CustomerId** which identifies the Customer Organization who the template belongs to. It also contains a **DocumentType**, specifying for which kind of document it is a template for. The last piece of information it contains is a **TimeStamp** specifying when the system received the template.